
TensorRT基础-概述

TensorRT基础

1. TensorRT的核心在于对模型算子的优化（**合并算子**、利用GPU特性**选择特定核函数**等多种策略），通过tensorRT，能够在Nvidia系列GPU上获得最好的性能
2. 因此tensorRT的模型，需要在目标GPU上**实际运行**的方式选择最优算法和配置
3. 也因此tensorRT生成的模型只能在**特定条件**下运行（编译的trt版本、cuda版本、编译时的GPU型号）
4. 主要知识点，是**模型结构定义方式、编译过程配置、推理过程实现、插件实现、onnx理解**


不跑跑，我哪知道怎么才最快

<https://www.cnblogs.com/qccz123456/p/11767858.html>


<https://www.bilibili.com/video/BV1Xw411f7FW/>

TensorRT基础

Un-Optimized Network



TensorRT Optimized Network



C++接口

TensorRT提供的C++接口

```
// Create input tensor of shape { 1, 1, 28, 28 }
ITensor* data = network->addInput(
    mParams.inputTensorNames[0].c_str(), DataType::kFLOAT, Dims3{1, mParams.inputH, mParams.inputW});
ASSERT(data);

// Create scale layer with default power/shift and specified scale parameter.
const float scaleParam = 0.0125f;
const Weights power{DataType::kFLOAT, nullptr, 0};
const Weights shift{DataType::kFLOAT, nullptr, 0};
const Weights scale{DataType::kFLOAT, &scaleParam, 1};
IScaleLayer* scale_1 = network->addScale(*data, ScaleMode::kUNIFORM, shift, scale, power);
ASSERT(scale_1);

// Add convolution layer with 20 outputs and a 5x5 filter.
IConvolutionLayer* conv1 = network->addConvolutionNd(
    *scale_1->getOutput(0), 20, Dims{2, {5, 5}}, mWeightMap["conv1filter"], mWeightMap["conv1bias"]);
ASSERT(conv1);
conv1->setStride(DimsHW{1, 1});

// Add max pooling layer with stride of 2x2 and kernel size of 2x2.
IPoolingLayer* pool1 = network->addPoolingNd(*conv1->getOutput(0), PoolingType::kMAX, Dims{2, {2, 2}});
ASSERT(pool1);
pool1->setStride(DimsHW{2, 2});

// Add second convolution layer with 50 outputs and a 5x5 filter.
IConvolutionLayer* conv2 = network->addConvolutionNd(
    *pool1->getOutput(0), 50, Dims{2, {5, 5}}, mWeightMap["conv2filter"], mWeightMap["conv2bias"]);
ASSERT(conv2);
conv2->setStride(DimsHW{1, 1});

// Add second max pooling layer with stride of 2x2 and kernel size of 2x3
IPoolingLayer* pool2 = network->addPoolingNd(*conv2->getOutput(0), PoolingType::kMAX, Dims{2, {2, 2}});
ASSERT(pool2);
pool2->setStride(DimsHW{2, 2});

// Add fully connected layer with 500 outputs.
IFullyConnectedLayer* ip1
    = network->addFullyConnected(*pool2->getOutput(0), 500, mWeightMap["ip1filter"], mWeightMap["ip1bias"]);
ASSERT(ip1);

// Add activation layer using the ReLU algorithm.
IActivationLayer* relu1 = network->addActivation(*ip1->getOutput(0), ActivationType::kRELU);
ASSERT(relu1);
```

Python接口

TensorRT提供的Python接口

```
def populate_network(network, weights):
    # Configure the network layers based on the weights provided.
    input_tensor = network.add_input(name=ModelData.INPUT_NAME, dtype=ModelData.DTYPE, shape=ModelData.INPUT_SHAPE)

    conv1_w = weights['conv1.weight'].numpy()
    conv1_b = weights['conv1.bias'].numpy()
    conv1 = network.add_convolution(input=input_tensor, num_output_maps=20, kernel_shape=(5, 5), kernel=conv1_w, bias=conv1_b)
    conv1.stride = (1, 1)

    pool1 = network.add_pooling(input=conv1.get_output(0), type=trt.PoolingType.MAX, window_size=(2, 2))
    pool1.stride = (2, 2)

    conv2_w = weights['conv2.weight'].numpy()
    conv2_b = weights['conv2.bias'].numpy()
    conv2 = network.add_convolution(pool1.get_output(0), 50, (5, 5), conv2_w, conv2_b)
    conv2.stride = (1, 1)

    pool2 = network.add_pooling(conv2.get_output(0), trt.PoolingType.MAX, (2, 2))
    pool2.stride = (2, 2)

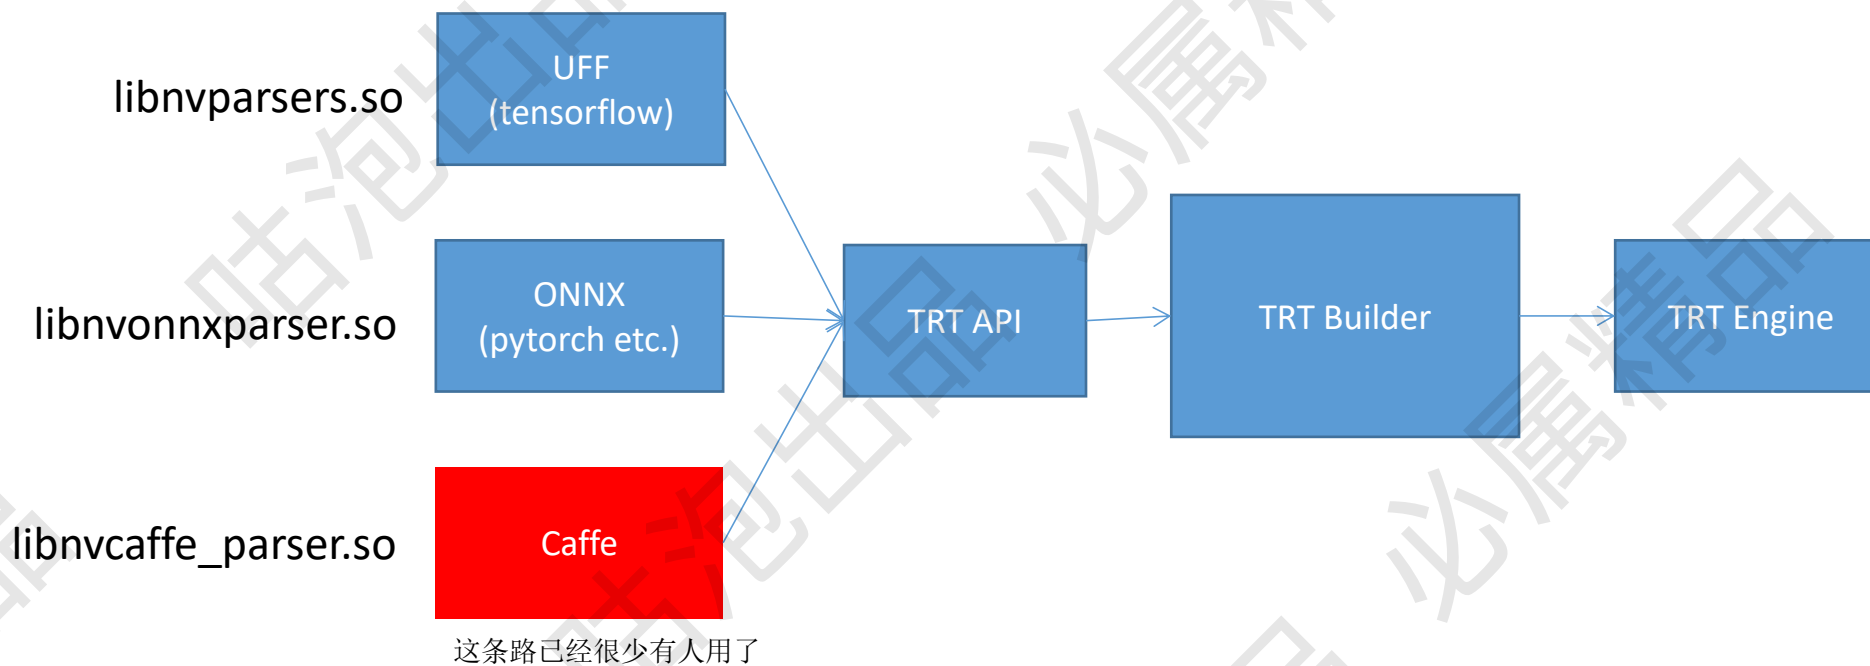
    fc1_w = weights['fc1.weight'].numpy()
    fc1_b = weights['fc1.bias'].numpy()
    fc1 = network.add_fully_connected(input=pool2.get_output(0), num_outputs=500, kernel=fc1_w, bias=fc1_b)

    relu1 = network.add_activation(input=fc1.get_output(0), type=trt.ActivationType.RELU)

    fc2_w = weights['fc2.weight'].numpy()
    fc2_b = weights['fc2.bias'].numpy()
    fc2 = network.add_fully_connected(relu1.get_output(0), ModelData.OUTPUT_SIZE, fc2_w, fc2_b)

    fc2.get_output(0).name = ModelData.OUTPUT_NAME
    network.mark_output(tensor=fc2.get_output(0))
```


工作流程



常见方案

基于tensorRT的发布，又有人在之上做了工作
<https://github.com/wang-xinyu/tensorrtx>

为每个模型写硬代码，并已写好了大量的常见模型代码




```
/* ----- yolov5 backbone----- */  
auto focus0 = focus(network, weightMap, *data, 3, get_width(64, gw), 3, "model.0");  
auto conv1 = convBlock(network, weightMap, *focus0->getOutput(0), get_width(128, gw), 3, 2,  
auto bottleneck_CSP2 = C3(network, weightMap, *conv1->getOutput(0), get_width(128, gw), get  
auto conv3 = convBlock(network, weightMap, *bottleneck_CSP2->getOutput(0), get_width(256, g  
auto bottleneck_csp4 = C3(network, weightMap, *conv3->getOutput(0), get_width(256, gw), get  
auto conv5 = convBlock(network, weightMap, *bottleneck_csp4->getOutput(0), get_width(512, g  
auto bottleneck_csp6 = C3(network, weightMap, *conv5->getOutput(0), get_width(512, gw), get  
auto conv7 = convBlock(network, weightMap, *bottleneck_csp6->getOutput(0), get_width(1024,  
auto spp8 = SPP(network, weightMap, *conv7->getOutput(0), get_width(1024, gw), get_width(10
```

方案思考



本课程主要学习以onnx路线的模型编译、推理和部署，原因主要有：

若使用onnx，则导出或者修改好的onnx模型，可以**轻易的移植到其他引擎上**、例如ncnn、rknn，这一点硬代码无法做到。并且用于排查错误，修改调整时也非常方便



获取代码

对应于系列名称: tensorrt-basic

获取代码: trtpy get-series tensorrt-basic

查询系列清单: trtpy series-detail tensorrt-basic

案例清单

```
C:\Users\Administrator\cuda-driver-api>trtpy series-detail tensorrt-basic
Use cache C:\Users\Administrator\.cache/trtpy\code_template\tensorrt-basic.series.json
List templ:
chapter: 1.1, caption: hello-tensorrt, description: 开始第一个tensorRT的旅程吧，编译一个模型
chapter: 1.2, caption: hello-inference, description: 编译好的模型进行推理
chapter: 1.3, caption: cnn-and-dynamic-shape, description: CNN结构的动态shape如何控制，要点在哪里
chapter: 1.4, caption: onnx-parser, description: 使用onnx解析器读取onnx文件构建模型结构和填充权重
chapter: 1.5, caption: onnx-parser-source-code, description: 使用onnx解析器的源代码，了解解析器细节原理
chapter: 1.6, caption: onnx-editor, description: 对onnx文件进行编辑、创建、读取
chapter: 1.7, caption: hello-plugin, description: 开始第一个插件，基于onnx的插件
chapter: 1.8, caption: integrate-easyplugin, description: 对插件做封装，插件开发更简单，更容易
chapter: 1.9, caption: int8, description: tensorRT的int8标量化
```

TensorRT的库文件一览

| | |
|---|---|
| <code>> stubs</code> | |
| <code>libnvcaffe_parser.a</code> | caffe的模型解析器, 输入caffemodel, 解析到tensorRT的模型结构INetworkDefinition |
| <code>libnvcaffe_parser.so</code> | |
| <code>libnvcaffe_parser.so.8</code> | |
| <code>libnvcaffe_parser.so.8.2.3</code> | |
| <code>libnvinfer_builder_resource.so.8.2.3</code> | |
| <code>libnvinfer_plugin_static.a</code> | nvidia提供的插件, 编译自这里的代码 https://github.com/NVIDIA/TensorRT/tree/main/plugin |
| <code>libnvinfer_plugin.so</code> | |
| <code>libnvinfer_plugin.so.8</code> | |
| <code>libnvinfer_plugin.so.8.2.3</code> | |
| <code>libnvinfer_static.a</code> | tensorRT的核心库 |
| <code>libnvinfer.so</code> | |
| <code>libnvinfer.so.8</code> | |
| <code>libnvinfer.so.8.2.3</code> | |
| <code>libnvonnxparser_static.a</code> | ONNX模型解析器, 输入onnx模型, 解析到tensorRT的模型结构INetworkDefinition |
| <code>libnvonnxparser.so</code> | |
| <code>libnvonnxparser.so.8</code> | |
| <code>libnvonnxparser.so.8.2.3</code> | |
| <code>libnvparsers_static.a</code> | uff模型解析器, 输入uff模型, 解析到tensorRT的模型结构INetworkDefinition |
| <code>libnvparsers.so</code> | |
| <code>libnvparsers.so.8</code> | |
| <code>libnvparsers.so.8.2.3</code> | |
| <code>libonnx_proto.a</code> | onnx的proto编译后的效果, 是由protoc产生的onnx.cc编译而成 |
| <code>libprotobuf-lite.a</code> | |
| <code>libprotobuf.a</code> | protobuf的静态库 |

谢谢!